

# Creative Machine Learning

*General introduction*

---

Pr. Philippe Esling  
[esling@ircam.fr](mailto:esling@ircam.fr)

---



# General introduction

---

## Artificial *intelligence* ?

- Requires to first understand what intelligence is !
- Must have something to do with *thinking*
- But in a broader sense, might also be about *perception* and *action*

## Why bothering at all ?

- In terms of philosophy, we would talk about problems involving these
- However, we *try to understand the mechanisms of our thinking itself*
- We seek *models* targeted at thinking, perception and actions

## Our overarching goal

Construct *mathematical models* reproducing *human behaviors*  
Hopefully helping us in understanding our own thinking process

# Artificial *intelligence* ?

---

Hard to define intelligence and thinking

Creation ex-nihilo does not exist  
(Laplace)

« *Man can only desire what he has already perceived* »  
**William Blake**

———— Try to invent a new animal right now ——

Intelligence: *the capacity to assemble two ideas that seemed heterogeneous*  
We exist (morphologically) since ~200 kYears, and around ~50 kYears

« *Take two concepts and create a third one without impairing the two first* »  
**Noam Chomsky**

So thinking is finding similarity, discriminating and seeing common patterns

*NB: Frames of Mind: the Theory of multiple intelligence – Howard Gardner*  
(logico-mathematic, spatial, inter-personal, corporal, linguistic, intra-personnal, musical, ecologist, existential)

To have a model, we will need a *representation* that exposes *constraints* on thinking

# Brief history of AI

**Engineer view**

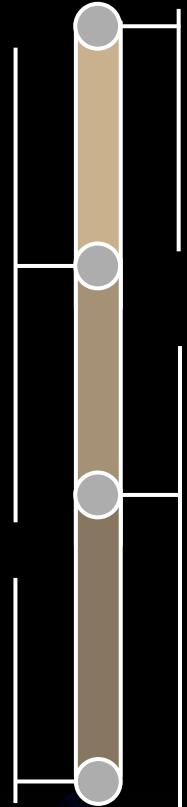
« Building smarter programs »

**Scientific secret goal**

« Build a computational account of intelligence »

**1950 - Alan Turing**

Inventor of *Bombe* and Turing test  
(code-breaking machine that turned the war)



**1842 - Lady Ada Lovelace**

First-ever computer programmer  
(decades before computers even existed)

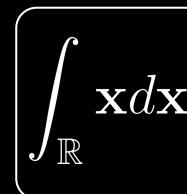
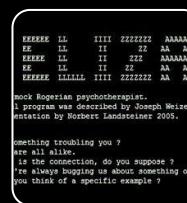
*The analytical engine has no pretensions to originate anything,  
it can do whatever we know how to order it to perform*

**1955 - Marvin Minski**

Coins the term *artificial intelligence*  
“Steps towards artificial intelligence”



**1966 - Shakey**  
First *electronic person* at Stanford  
(general mobile robot with reasoning)



**1960s - Early dawn age**

**1966 - ELIZA Chatbot**

First-ever chatbot (*Rogerian psychotherapy*)  
Beginning of Natural Language Processing (NLP)

<https://www.masswerk.at/elizabot/>

**1967 - Symbolic integration**

Program written by Joel Moses at MIT  
Scores better than 98% of the students

**1970 - 1980**



**1st AI Winter**

# Brief history of AI

## 1943 - Neuron

First model by McCulloch & Pitts (purely theoretical)



## 1957 - Perceptron

Actual **learning machine** built by Frank Rosenblatt

Learns character recognition analogically



## Lesson #2 1986 - Backpropagation

First to learn neural networks efficiently (G. Hinton)



## Lesson #3 1989 - Convolutional NN

Mimicking the vision system in cats (Y. LeCun)



## Lesson #4 2012 - Deep learning

Layerwise training to have deeper architectures

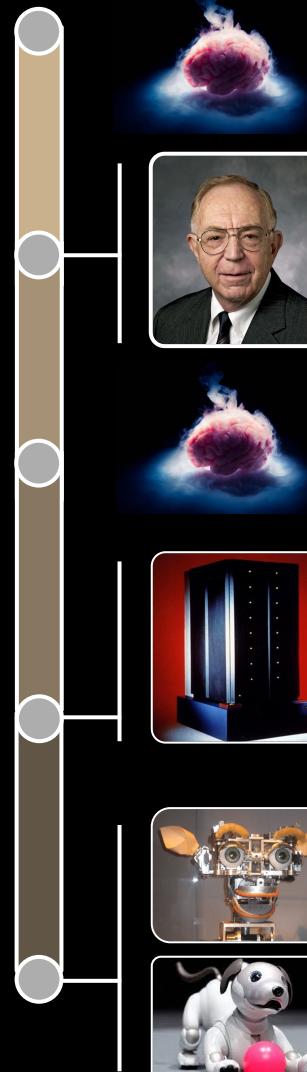
Swoop all state-of-art in classification competitions



## Lesson #6 2015 - Generative model

First wave of interest in generating data

Led to current model craze (VAEs, GANs, Diffusion)



## 1970 - 1980 1st AI Winter

## 1980s - Expert systems

Emulates decisions of human experts

Rising use in healthcare (and still today)

## 1985 - 1995 2nd AI Winter

## 1997 - Deep blue

First defeat of human chess champion

Based on alpha-beta min-max algorithm

## 2000s - Robotic explosion

### 1998 - Kismet

First emotional robot

### 1999 - Aibo

Robotic dog

### 2002 - Roomba

Robotic vacuum

### 2011 - Siri

Personal assistant

2012 onwards Deep learning era

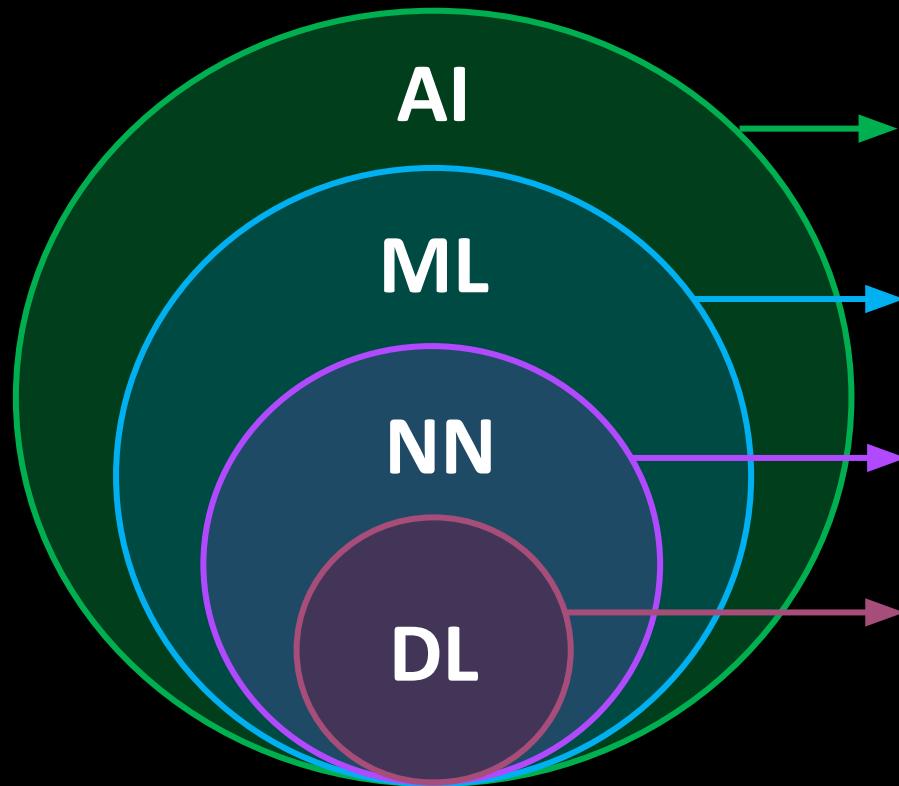
# Families of Artificial Intelligence

## Solving confusion around machine *intelligence*

There is wide controversies about AI (*weak vs. strong, feasible vs. fake*)

Avoiding any debate, we will only discuss here *machine learning*.

- Can we automatically learn parameters of an approximation
- The most active field recently, and also very wide topics inside it



### **Artificial Intelligence (AI)**

*Any technique allowing machines to solve human tasks*

### **Machine Learning (ML)**

*Learning inference models from examples*

### **Neural Networks (NN)**

*Brain-inspired ML models*

### **Deep Learning (DL)**

*Building (deep) hierarchies of NN representations*

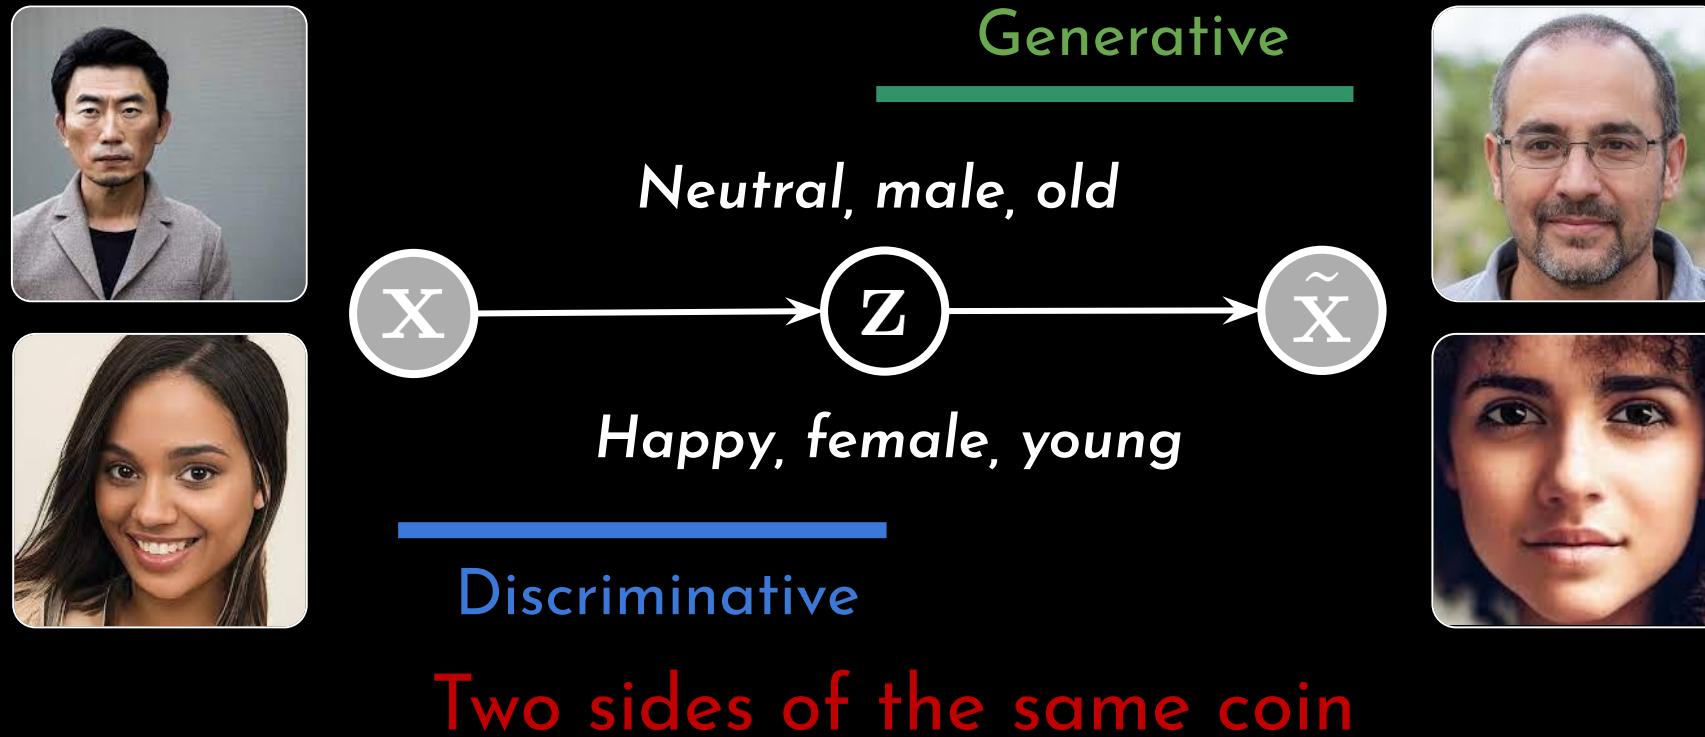
# Why creative ?

## 3 major types of learning

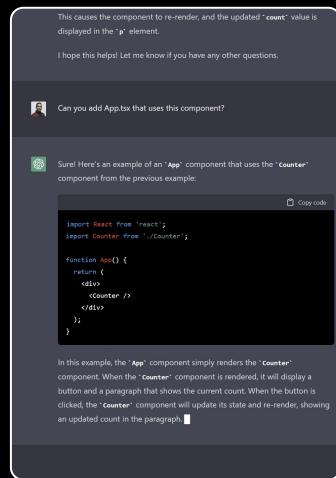
- **Supervised learning** is inferring a function from labeled training data
- **Unsupervised learning** is trying to find hidden structure in unlabeled data
- **Reinforcement learning** is acting to maximize a notion of cumulative reward

**Self-supervised learning** is using the data itself directly to define tasks to solve

## Discriminative or generative ?



# Major players in the generative field



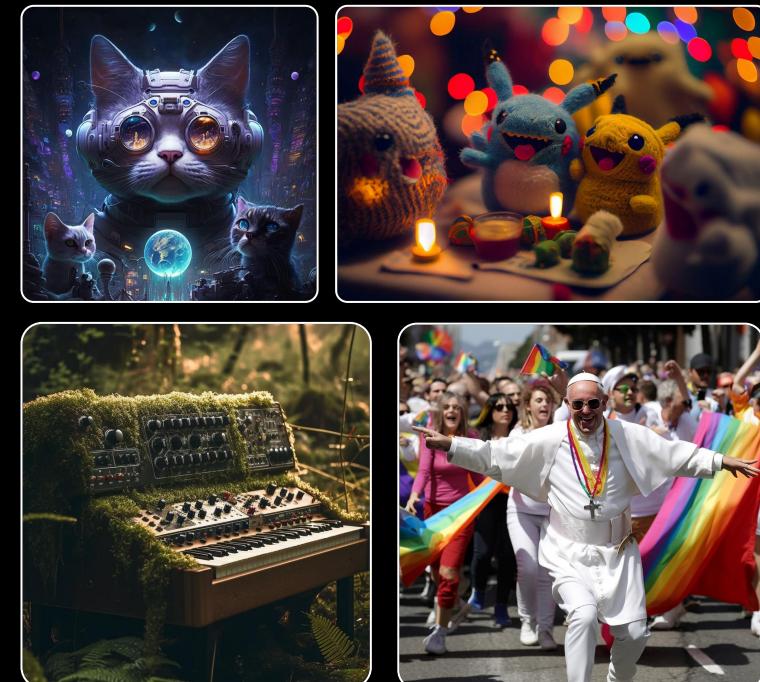
## ChatGPT

AI language model (OpenAI) based on Transformers (GPT)

Trained on vast corpus of text to generate chat answers

Impressive question answering and reasoning abilities

### *Lesson #5 = Transformers*



## Midjourney - StableDiffusion



Image-generation AI by Stability based on GANs and diffusion

Trained to generate HQ images based on prompts

Impressive capabilities used in art, entertainment, and fashion

### *Lesson #9 = Generative Adversarial Networks (GANs)*

### *Lesson #10 = Diffusion models*

# What about musical data ?

---



*Creep by Radiohead  
... but by The Beatles*



Generative model by OpenAI creates original music with lyrics and singing.

*Combination of VQ-VAE and Transformer for raw audio at 44.1 kHz.*

*Can generate music in various styles, and even mimic specific artists*

MusicLM



*A fusion of reggaeton and electronic dance music, with a spacey, otherworldly sound.  
Induces the experience of being lost in space [...]*



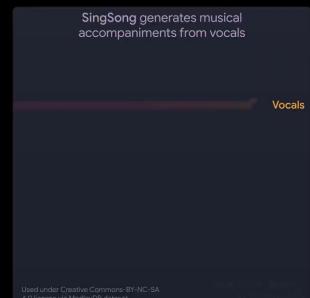
## Voice cloning

Text-to-speech allows to clone any voice  
Cyberduck app

*Thank you Barack !*

## Reverse Karaoke

Generate music to accompany vocals  
SingSong project



Used under Creative Commons BY-NC-SA 4.0 license via MedleyDB dataset

# Course program

---

## Goal of this course

- Ability to construct (math) and develop (code) your own *generative models*
- True and deep understanding of the mathematical theory behind those models
- Constant flow between theory, implementation and creative application

## Detailed program

Discriminative

Generative

1. Introduction
2. Machine learning
3. Neural networks
4. Advanced networks
5. Deep learning
6. Probabilities and Bayesian inference
7. Latent models, EM and GMM
8. Approximate inference
9. Variational Auto-Encoders (VAEs) and flows
10. Adversarial learning (GANs)
11. Diffusion models

# Why should you even listen to me

---



**Professor** Philippe Esling - [esling@ircam.fr](mailto:esling@ircam.fr)



**Sorbonne université (Paris)**

Teaching machine learning for 10 years



**The University of Tokyo**

Invited professor in ML



ircam



ACIDS

**ACIDS - Artificial Creative Intelligence and Data Science**

Lead of research group (in A/S team) on generative models for music



**Japanese-French Laboratory for Informatics**

Invited researcher for leading the Creative AI axis



**Google Research (Brain)**

Invited researcher at the Magenta team

Presentation of my research in the second half of this course

---

## Guest lecturers



**Taketo Akama**

Industrial research  
Sony CSL Tokyo  
Generative models



**Pr. Tatsuya Harada**

Academic research  
University of Tokyo  
Computer vision, robotics

# Course information

---

## Basic administrative information

**Type** : 2-credits graduate school course

**Period** : April - July 2023

**Span** : 14 classes of 105 minutes

**Date** : Every thursday at 2:55 - 4:40 pm (JST)

**Onsite** : *Room 214, 2nd Floor, Sci. 7 Building*

### Online course

[https://u-tokyo-ac-jp.zoom.us/j/  
81363691008?pwd=cmxiRURMdM9udXBKbTNjQkZvbINFQT09](https://u-tokyo-ac-jp.zoom.us/j/81363691008?pwd=cmxiRURMdM9udXBKbTNjQkZvbINFQT09)



### Github page

[github.com/esling/creative\\_ml](https://github.com/esling/creative_ml)

### Professor

Philippe Esling - [esling@ircam.fr](mailto:esling@ircam.fr)

# Pre-requisites

---

## Notions required but not covered in this course

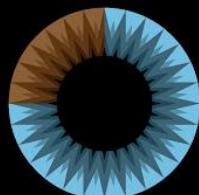
Resources to help you out if you feel behind on these topics

### Calculus

Limits and infinitesimals  
Differential calculus  
Integral calculus

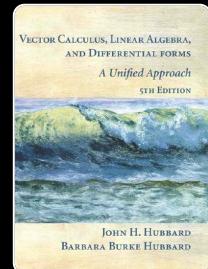
### Linear algebra

Vector spaces  
Matrix calculus  
Geometry and analysis



### 3 Blue 1 Brown

[Essence of linear algebra](#)  
[Essence of calculus](#)



### Hubbard J. & Hubbard B.B

*Vector Calculus, Linear Algebra, and Differential Forms: A Unified Approach* (5th edition)

- Linear algebra [slides](#)
- Probability [revision slides](#)
- Statistics [course notes](#)
- Sampling [pages](#)
- [Matrix Cookbook](#)

### Programming

Python basics  
Functional notions  
Object-oriented prog.

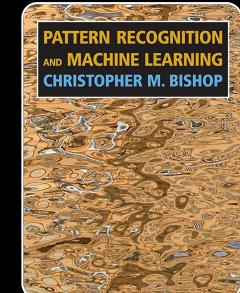
### Data analysis

Data visualization  
Basic filtering, processing



### Machine learning bible

**Christopher M. Bishop**  
*Pattern recognition and machine learning*



# Course teaser

## Overarching construction

Course heavily based on **interactive learning** principles

Merging **theory and application** as much as possible



#2 - Jupyter notebooks



#3 - Interactive panels

- Interactive coding environment with exercises
- Used during the class with dynamic panels
- Will also serve to fill in evaluated exercises
- Submission system impending

Fully online

Slides  
[slides.google.com](https://slides.google.com)



Colab  
[colab.google.com](https://colab.google.com)



GitHub **fork** (+ star)  
**Jupyterlab** environment

Best solution



**Course and exercises are always evolving**

**Help each others out** and work as groups

**Post issues** on GitHub and Discord

**Join discord**

<https://discord.gg/gpF82wnh>



# Understanding philosophy

## General principle

We will be exploring all concepts by

1. Deriving the mathematical bases ourselves
2. Implementing low-level aspects directly as well
3. Moving up to higher-level (simpler) implementation

Essential principle for a profound understanding of concepts

## ML Libraries

Concept mirrored in our choice of libraries

1. **Numpy** forces us to do mathematical derivation
2. **JAX** adds automatic differentiation over Numpy
3. **Pytorch** allows to define networks in seconds



**Numpy**



**JAX**



**PyTorch**

## Other helpful libraries



Librosa



Scipy



Seaborn



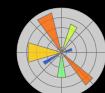
Bokeh



Scikit



Panel



Matplot



Manim

# Evaluation methodology

---

## One major evaluation and project



### Continuous evaluation

Course based on interactive Jupyter **notebooks to fill and submit**

- Each notebook contains a set of exercises
- We will start exercises together but you have to finish at home
- Automatic submission system will be set in place
- Rolling deadlines of 3 sessions (~3 weeks) to submit your work

### Project-based evaluation

**Supplementary evaluation** will be based on a large-scale project

- Several projects will be proposed
- Groups of **4 to 8 students** per project
- Projects are coded in **Python** (cf. Coding style section underneath)
- All projects should be **accompanied with a small report** in english
- The **report should be of 8 pages** following scientific papers style
- Reports are written in LaTeX with famous conferences (*ICLR, NeurIPS*) format style

# Setup and course follow-up

---

## Following the course

All lessons will follow the same structure

Mixture of animated slides and notebooks

You can follow the online or offline version

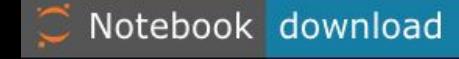
Slides

Notebooks

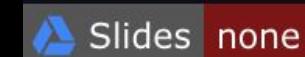
**Online**



**Offline**



**The course is updated prior to each lesson** Course not available yet



## Installing the course [github.com/esling/creative\\_ml](https://github.com/esling/creative_ml)

**You need a GitHub account**



1. **Star the repo** to get updates on new courses and corrections
2. **Fork the repo** to have your own copy (for exercises)
3. **Clone the repo** to your local machine

## Setup your environment



1. **Install jupyter lab** (see instructions in Setup section)
2. **Install all required libraries** (see the requirements.txt file)
3. **Test your setup** by using the 00\_setup.ipynb notebook

# [Setup time]

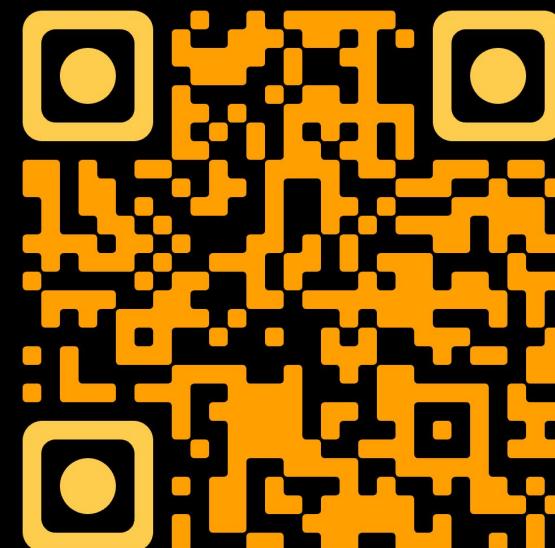
*5 minutes + 5 minutes break*

Github page



[github.com/esling/creative\\_ml](https://github.com/esling/creative_ml)

Installation guide



[creative\\_ml#setup](https://github.com/esling/creative_ml#setup)

# Machine learning

---

## Problem statement

We witness a phenomenon

Through a set of observations

$(x_1, y_1) \ (x_2, y_2) \ \dots \ (x_n, y_n)$

We know there exists a relation  
between  $\mathcal{X}$  and  $\mathcal{Y}$

But we do not know its nature

$$f : \mathcal{X} \rightarrow \mathcal{Y}$$

Machine learning constructs  
approximations

$$f_\theta : \mathcal{X} \rightarrow \mathcal{Y}$$

such that

$$y_i \approx \bar{y}_i = f_\theta(x_i)$$

# Regression and classification

---

The two most classical machine learning problems

Navigation icons: back, forward, search, etc.

# Machine learning

---

Formal problem statement

We aim to model the relationship between  $\mathcal{X}$  and  $\mathcal{Y}$

$$f : \mathcal{X} \rightarrow \mathcal{Y}$$

We collect a **dataset** that is representative of that relation

$$\mathcal{D} = \{(\mathbf{x}_1, \mathbf{y}_1) \cdots (\mathbf{x}_n, \mathbf{y}_n)\}$$

We will construct a **parametric approximation**  $f_\theta \in \mathcal{F}_\Theta$

$$f_\theta : \mathcal{X} \rightarrow \mathcal{Y} \quad \bar{\mathbf{y}} = f_\theta(\mathbf{x})$$

We want to find the best  $f_\theta^*$  so that  $\bar{\mathbf{y}} \approx \mathbf{y}$

So we need to find the best parameters  $\theta^* \in \Theta$

Need to compute the errors (**loss**) of our model

$$\mathcal{L}(\bar{\mathbf{y}}, \mathbf{y} \mid f_\theta, \theta)$$

And minimize (**optimize**) this amount of errors

$$\theta^* = \operatorname{argmin}_\theta \mathcal{L}(\bar{\mathbf{y}}, \mathbf{y} \mid f_\theta, \theta)$$

# Learning as error minimization

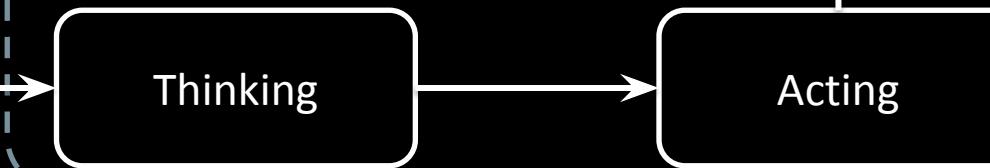
Why it makes sense in real-life

**Optimization**  
(error-minimization)



Feedback

Failing



- **Goal-directed learning** = acquiring, modifying or reinforcing knowledge
- Most problems can be summarized by a **trial-error process**
- Is a nickname for **optimization** (through **error minimization**)
  - So what can we do with ML ? Some examples from our team (ACIDS)

# ACIDS Project

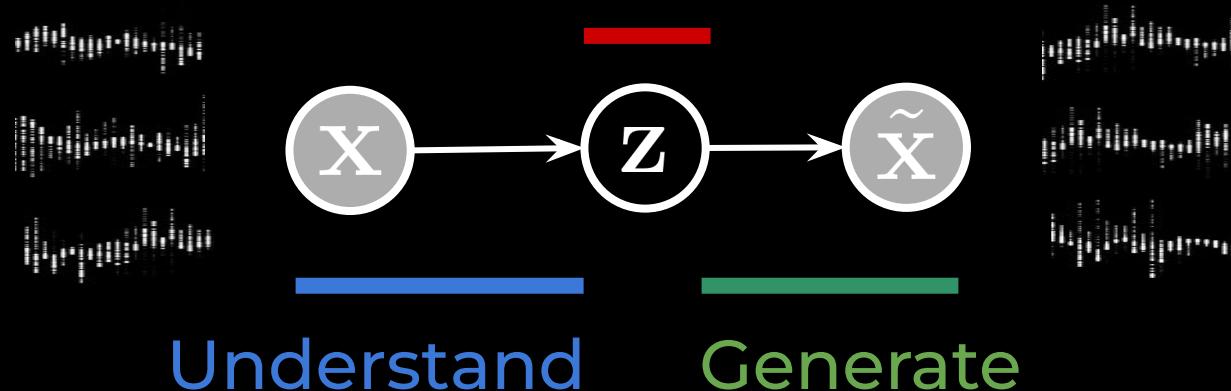


Artificial Creative Intelligence and Data Science

Research group at IRCAM, Sorbonne, Paris - Analysis / Synthesis team  
Academic machine learning research since 2012

Theoretical and applied research in generative models for musical creation

Control

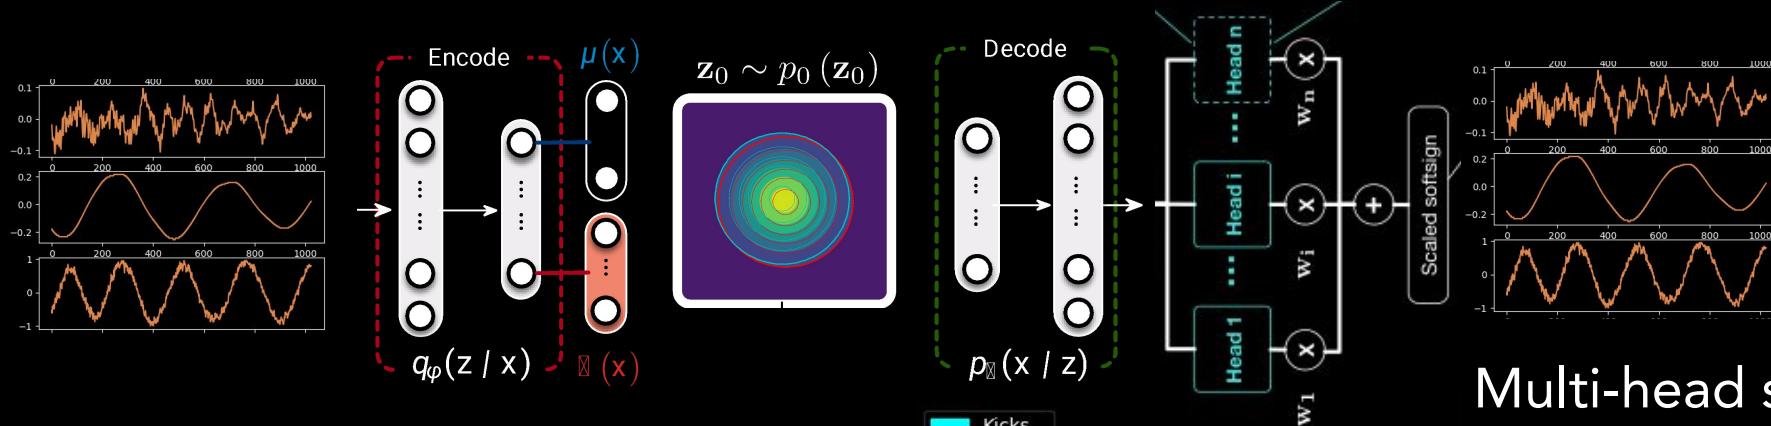


- Neural drum machine
- Timbre transfer applications
- High-level controls
- **[RAVE]** - Real-time synthesis
- **[FlowSynth]** - Synthesizer control
- **[Neurorack]** - Neural synthesizer

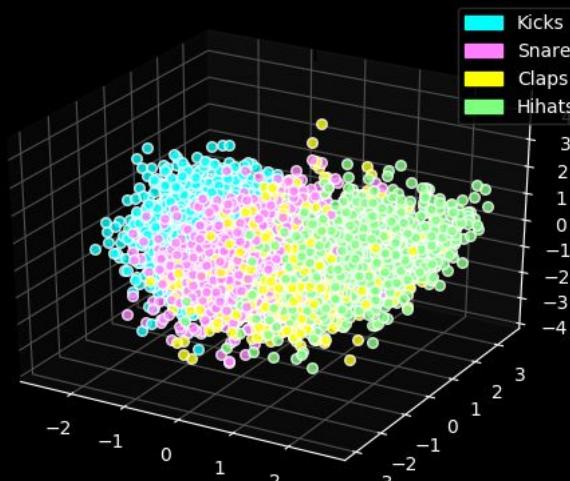
***Pushing the limits of musical creativity through the use of machine learning***  
***Creating the next generation of musical instruments***

# Applications - Drums generation

Drums are mostly noisy, require end-to-end training with waveforms

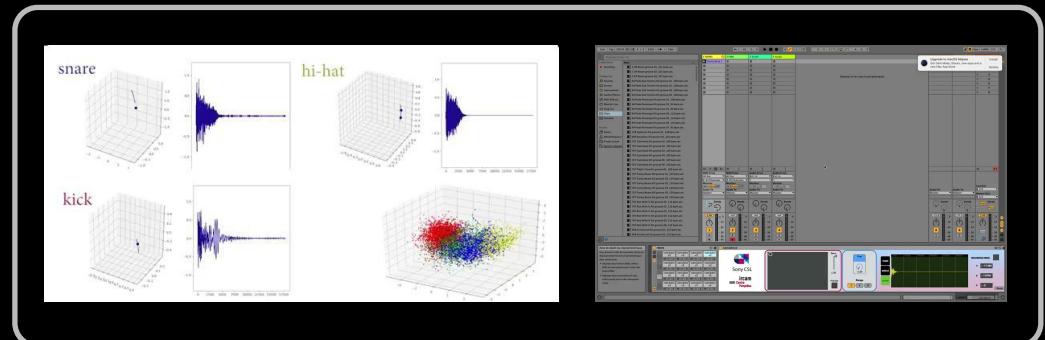


Work with Axel CRS and  
Cyran Aouameur



Real-time version demonstration

Multi-head spectrum inversion  
Heads unsupervisedly split into bands



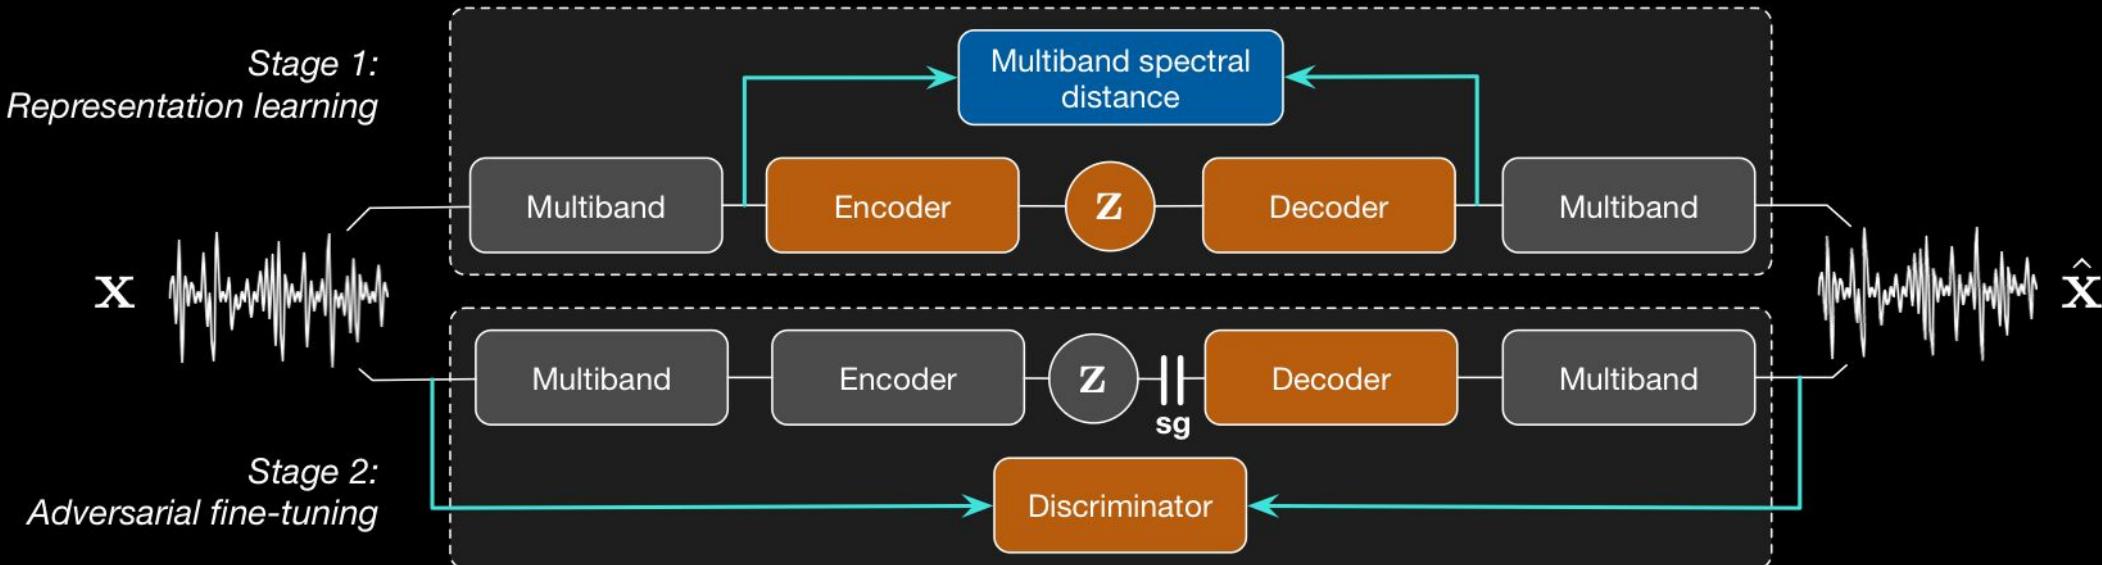
Ableton plugin and web-based interface

Aouameur, C., Esling, P. and Hadjeres G. Neural drum machine: An interactive system for real-time synthesis of drum sounds. arXiv:1907.02637, 2019.

# RAVE - Realtime audio waveform VAE

Recently, we vastly improved VAE generation *Full paper submitted to ICLR 2022*

Work with Antoine Caillon



**High-quality (48kHz) synthesis with 25x realtime on a laptop CPU**

Perform *timbre transfer* without any specific training →

Allows *polyphonic unconditional generation* of signals

Caillon, A. Esling, P. RAVE: A variational auto-encoder for fast and high-quality neural audio synthesis. arXiv:2111.05011, 2021.

# RAVE - Multiple outputs / interactions



**Realtime demo**



**Jetson embedding**

RAVE is available as multiple types of open source possibilities

<https://github.com/acids-ircam/RAVE>

PureData / MaxMSP

[https://github.com/acids-ircam/nn\\_tilde](https://github.com/acids-ircam/nn_tilde)

VST (multi-platform)

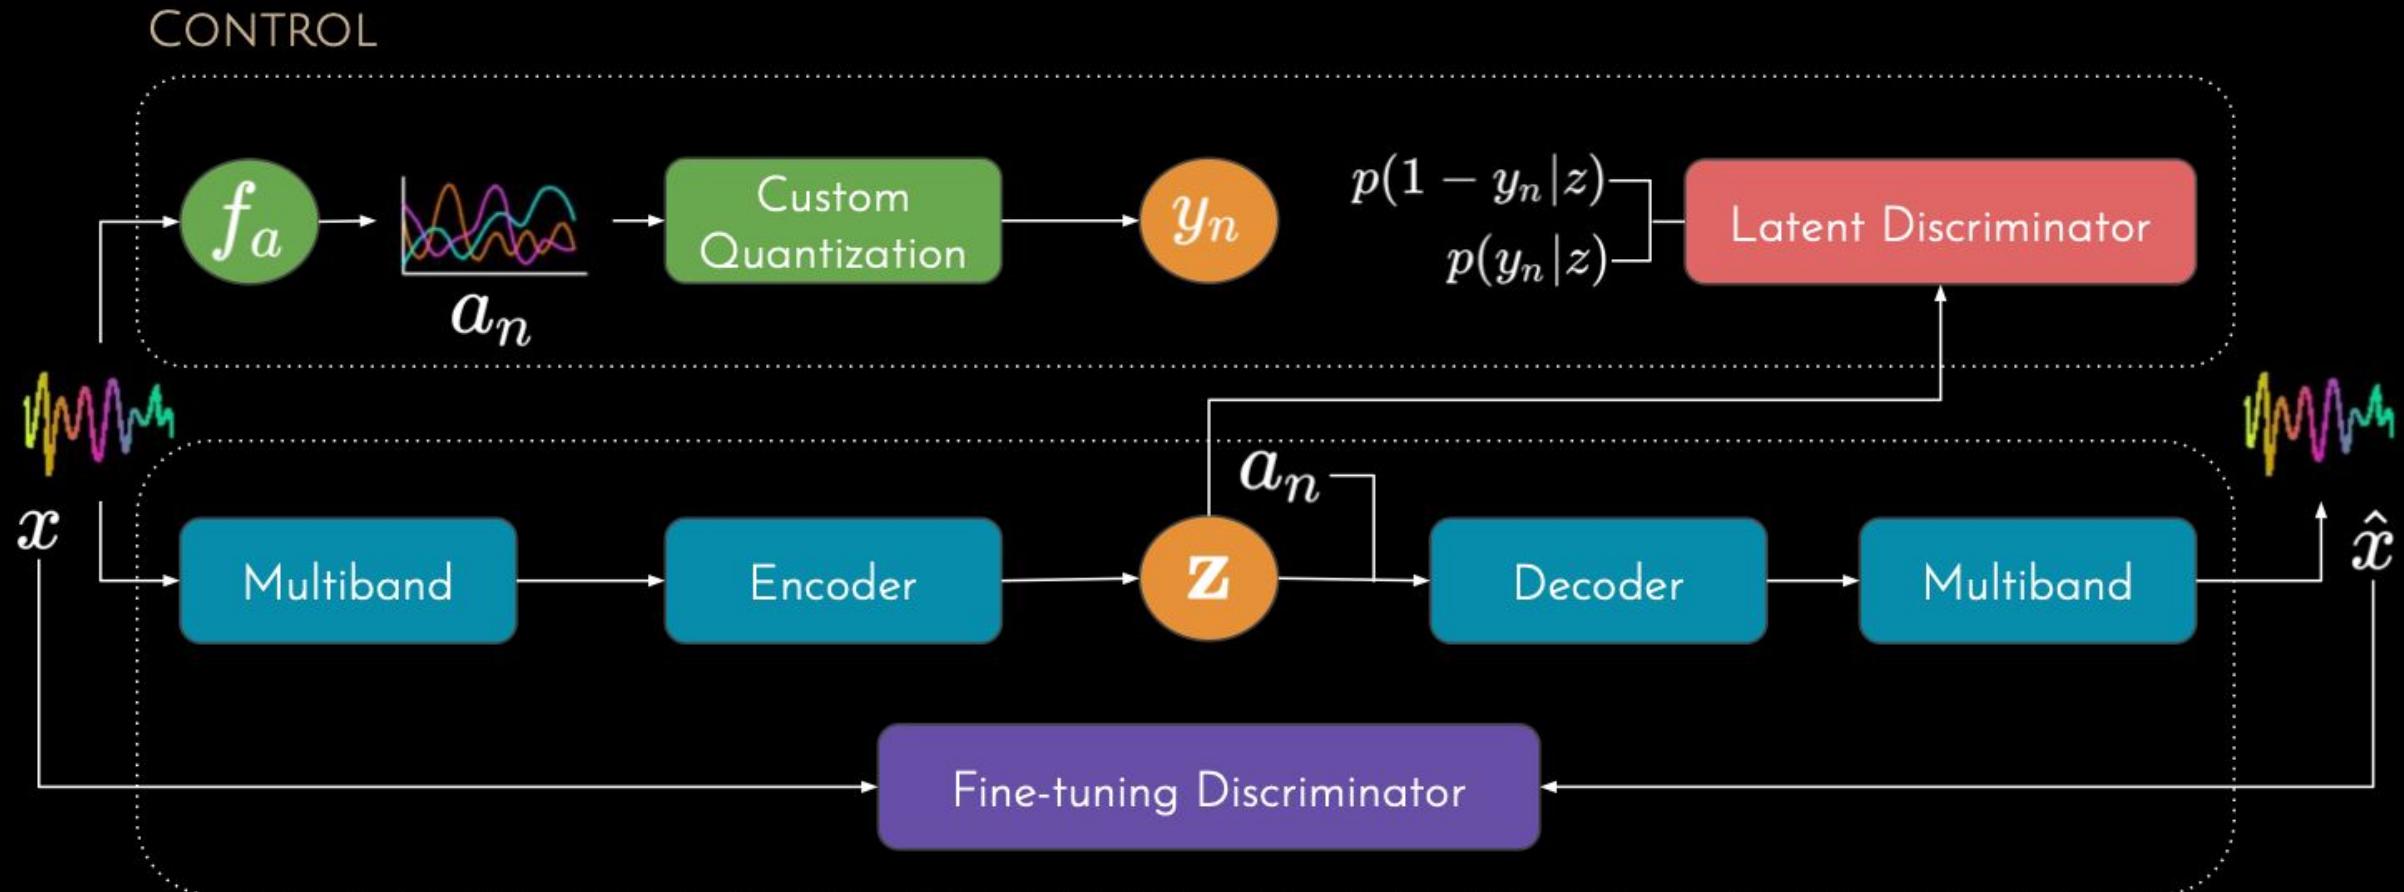
[https://github.com/acids-ircam/rave\\_vst](https://github.com/acids-ircam/rave_vst)

Caillon, A. Esling, P. RAVE: A variational auto-encoder for fast and high-quality neural audio synthesis. arXiv:2111.05011, 2021.

# Descriptor-based synthesis

Further improve the real-time generative model with high-level control

Work with Ninon Devis, Nils Demerlé, Sarah Nabi and David Genova



Devis, N. Demerle, N. Nabi, S. Genova, D. and Esling, P. Continuous descriptor-based control for deep audio synthesis. ICASSP 2023

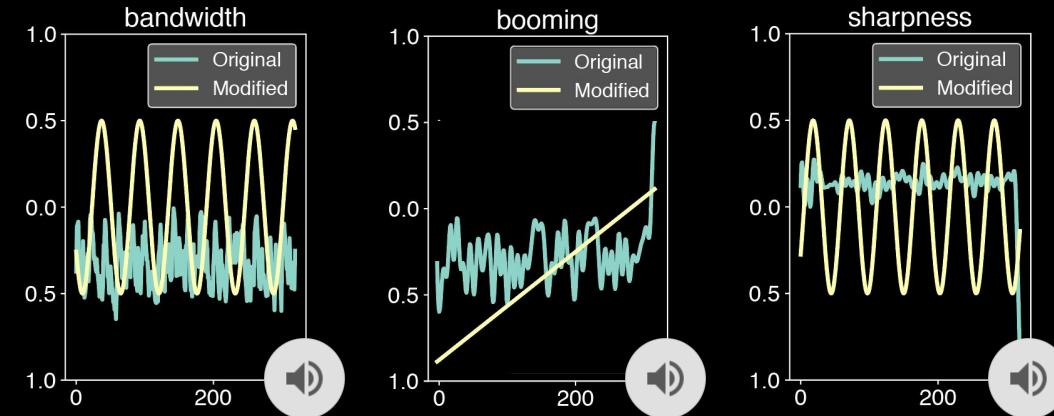
# Descriptor-based synthesis

## Descriptor-based synthesis (attribute transfer)

Model generation



(Derbouka)



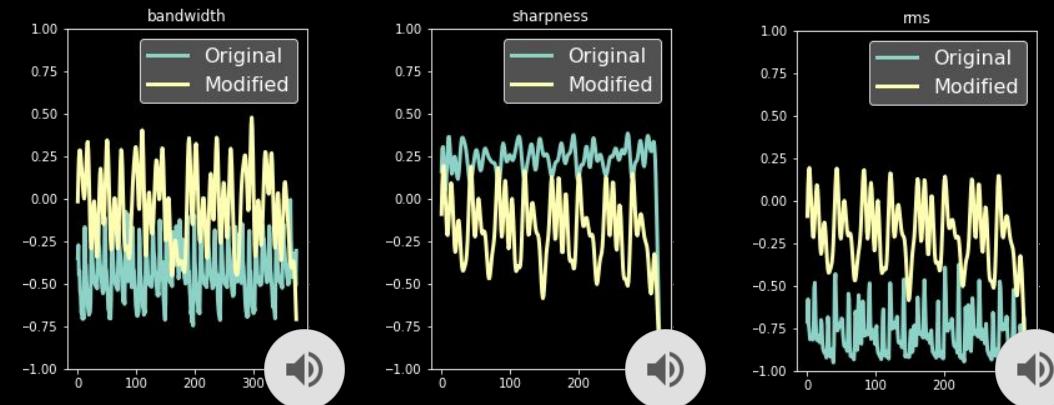
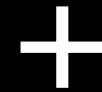
## Combining timbre and attribute transfer



(Derbouka)



(Japanese)



Devis, N. Demerle, N. Nabi, S. Genova, D. and Esling, P. Continuous descriptor-based control for deep audio synthesis. ICASSP 2023

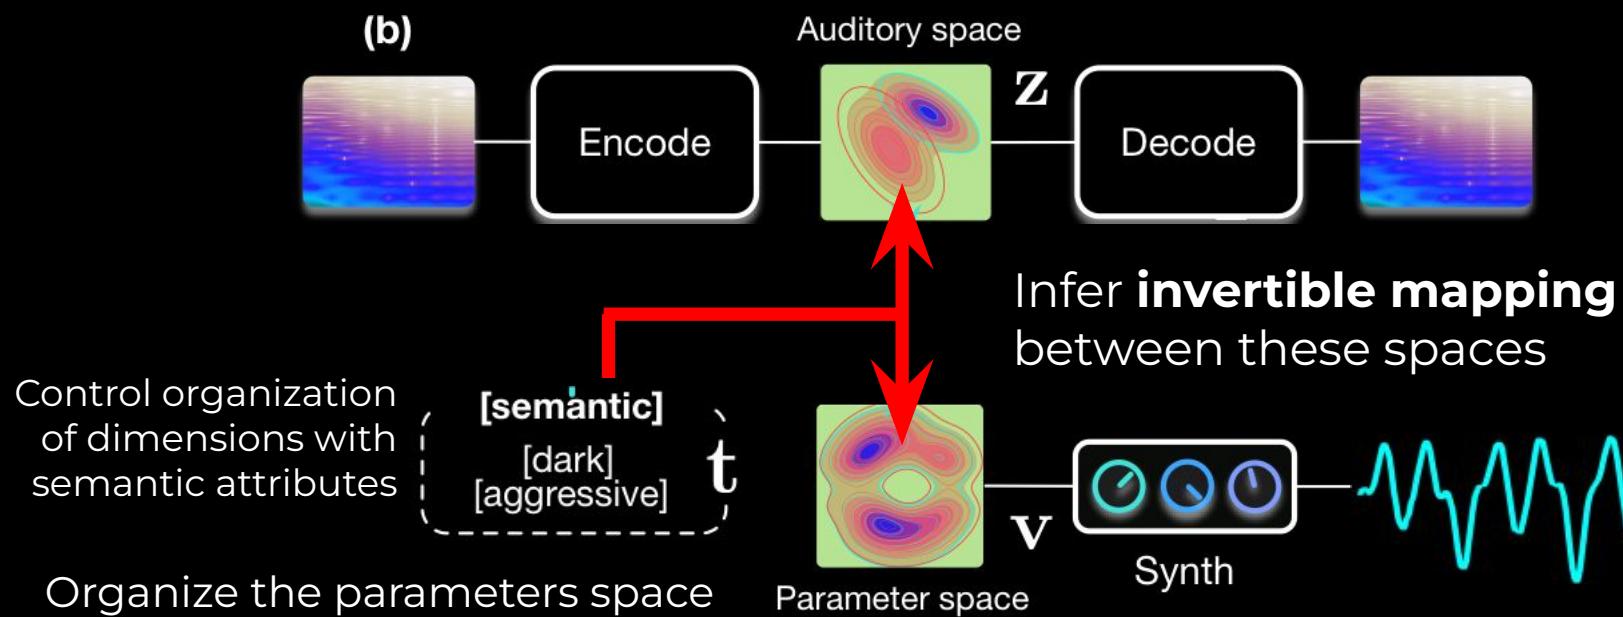
# Universal synthesizer control with FlowSynth

Compared to previous approaches (limited to parameters inference)



Our proposal = Universal synthesizer control with flows

Organize the auditory space of a synthesizer audio possibilities



# FlowSynth - Demo video

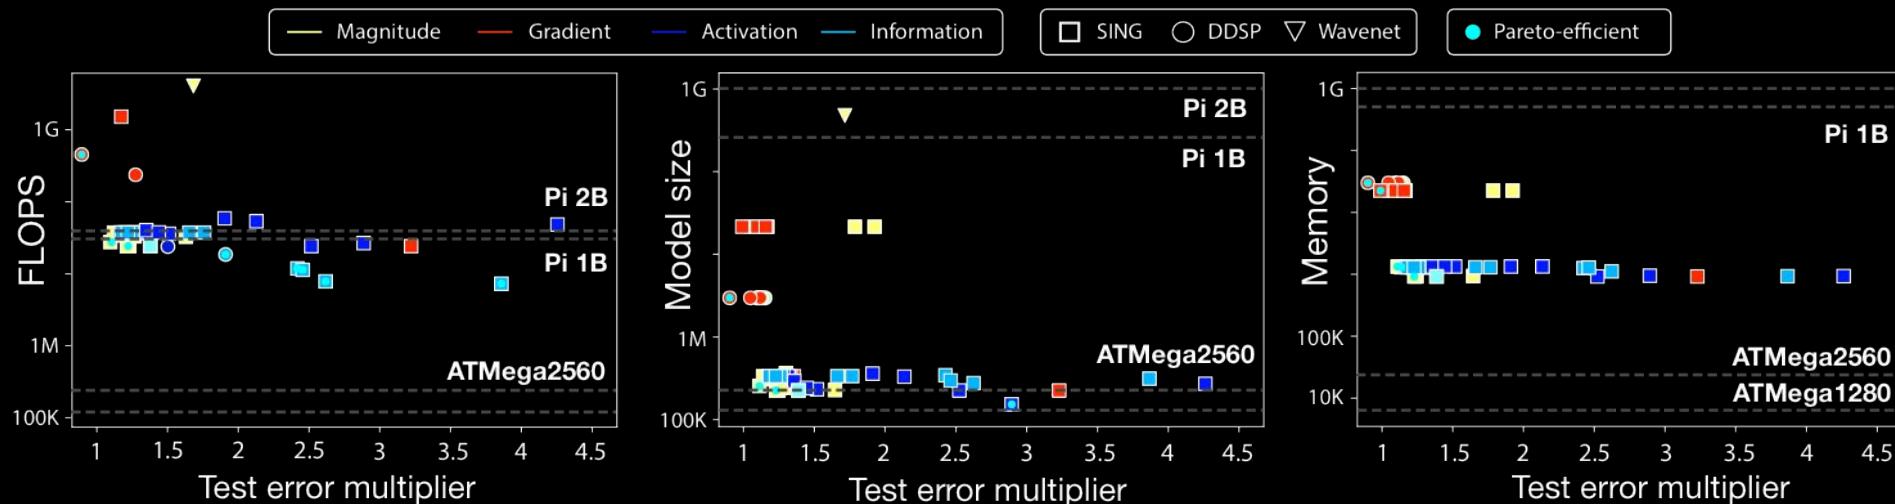
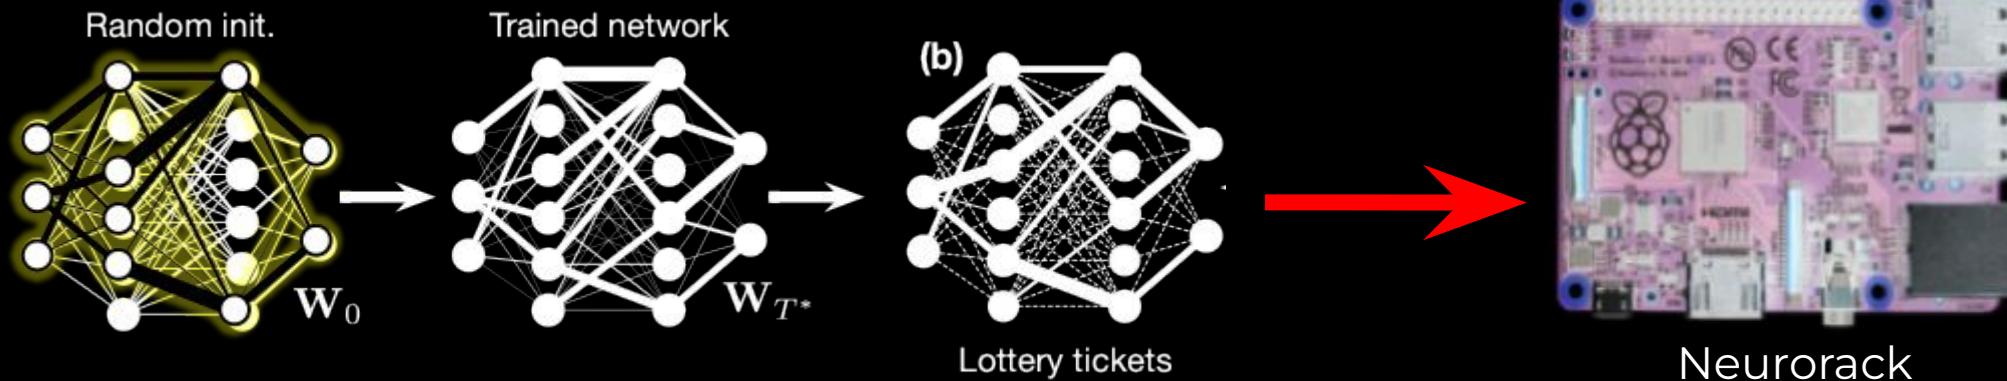


[https://github.com/acids-ircam/flow\\_synthesizer](https://github.com/acids-ircam/flow_synthesizer)

<https://www.youtube.com/watch?v=UufQwUitBlw>

# Lottery ticket hypothesis

Reducing deep models onto embedded chips

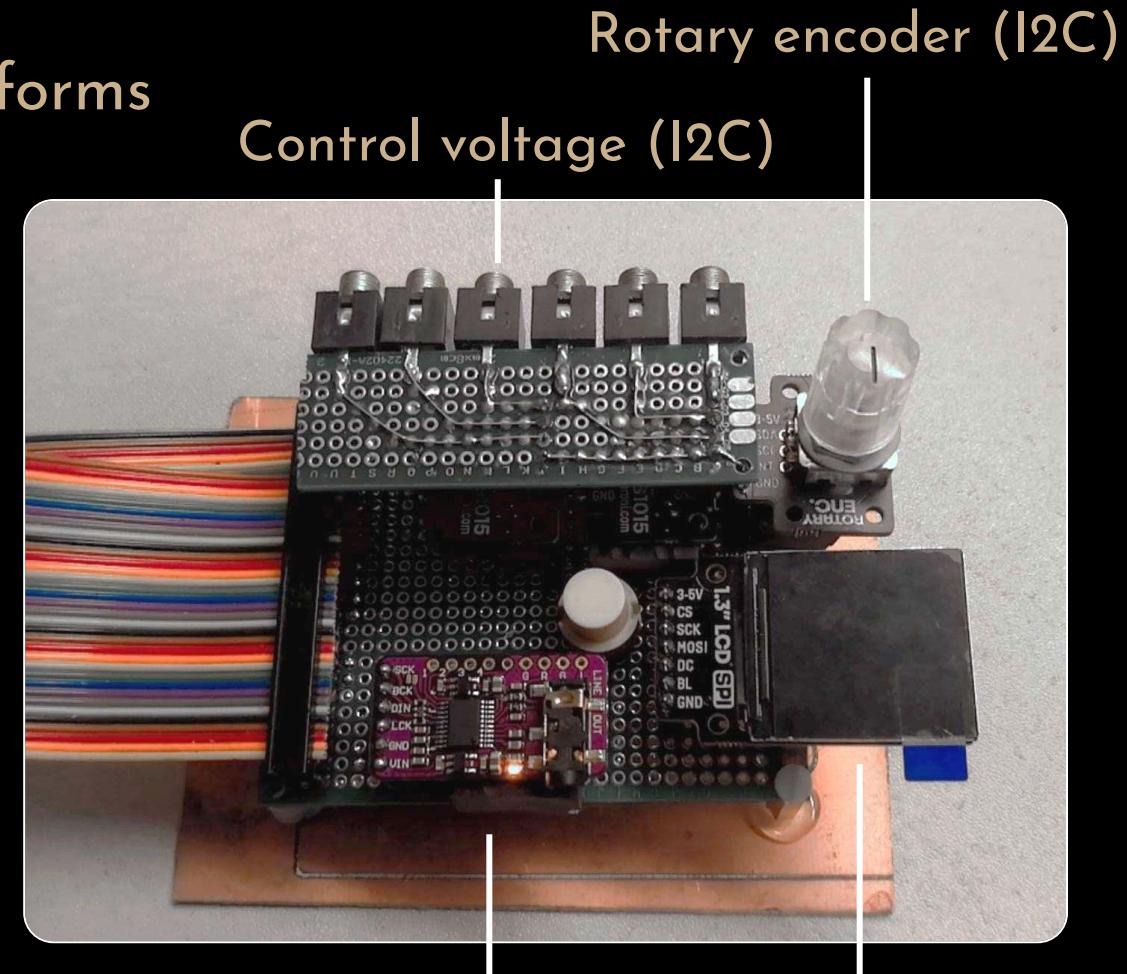
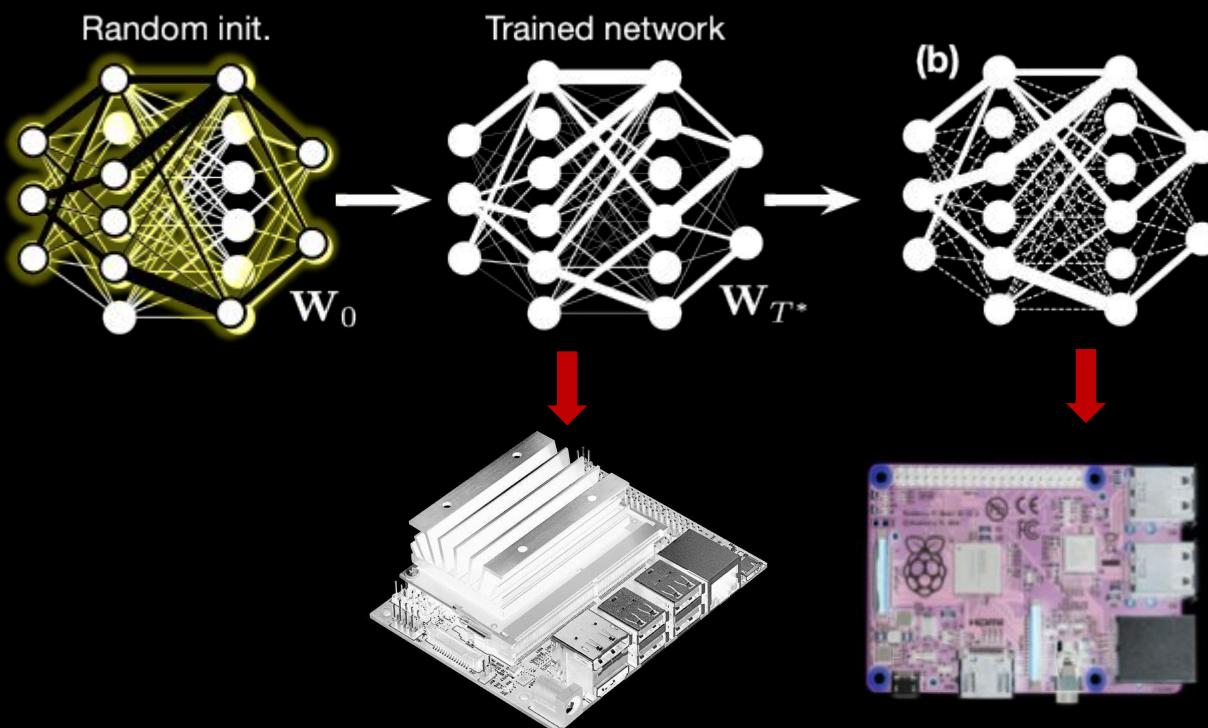


Esling, et al. *Diet deep generative audio models with structured lottery*. Proceedings of International Conference on Digital Audio Effects (DAFx), 2020

# Neurorack first prototype

Deep models in embedded constrained platforms

Neurorack // Work with **Ninon Devis** and **Martin Vert**

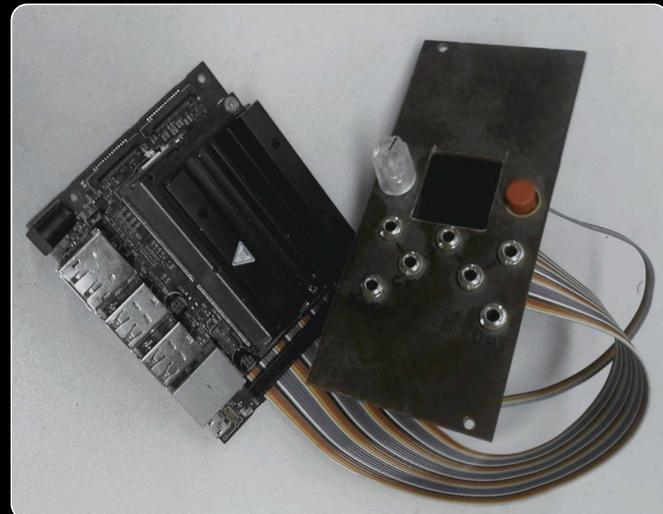


**First-ever Neurorack prototype built ourselves with simple components**

Devis, N. and Esling, P. *Neurorack: deep audio learning in hardware synthesizers*. EPFL PhD Conference (Lausanne), 2021

# Final neurorack prototype interface

Navigate the menu    Display interactive menu  
**Rotary**    **Screen**



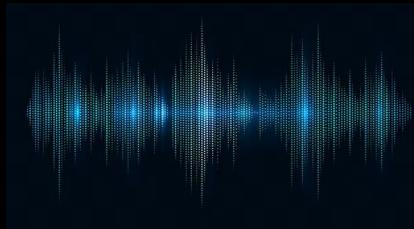
**Control Voltage**  
Vary sound features

**Gate**  
Trigger generation or interpolation

**Button**  
Select and click

**Audio out**

# Other applications (*discriminative*)



**Genre:** Acid techno  
**Tempo:** 140 BPM  
**Key:** B#  
**Lyrics:** Avin' it



Separating individual instruments or voices from a mixed audio signal.

Usual problem of frequency coverage and bandwidth

Recent models provide stunning results

Jansonn A. et al. (2017). "Singing Voice Separation with Deep U-Net Convolutional Networks." ISMIR.

## Music Information Retrieval (MIR)

### Music Genre Classification

Automatically classifying music genres from audio samples.

Nam, J et a. (2018). Deep learning for audio-based music classification and tagging. IEEE signal processing magazine, 36(1), 41-51  
[https://e-tarjome.com/storage/panel/fileuploads/2019-07-04/1562228921\\_E11422-e-tarjome.pdf](https://e-tarjome.com/storage/panel/fileuploads/2019-07-04/1562228921_E11422-e-tarjome.pdf)

### Beat and Tempo Detection

Analyzing music to detect beat patterns and tempo using deep learning.

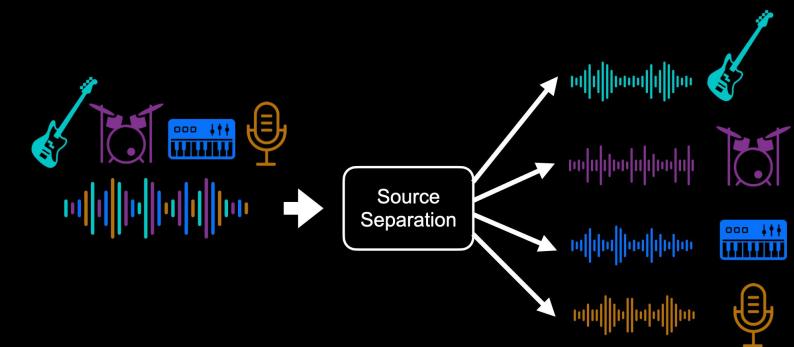
Schreiber, H., & Müller, M. (2018). "A Single-Step Approach to Musical Tempo Estimation Using a Convolutional Neural Network." ISMIR.  
[https://www.tagtraum.com/download/2018\\_schreiber\\_tempo\\_cnn.pdf](https://www.tagtraum.com/download/2018_schreiber_tempo_cnn.pdf)

### Automatic Music Tagging

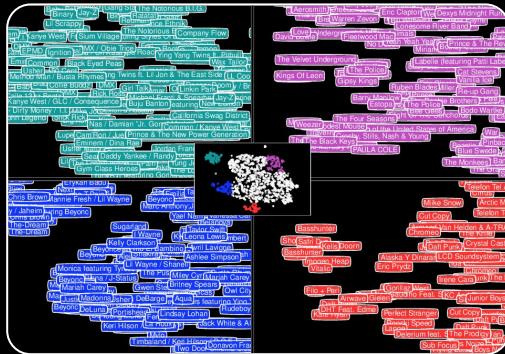
Assigning genre, mood, or other descriptive tags to songs using deep learning.

Choi, K., Fazekas, G., Sandler, M., & Cho, K. (2016). "Convolutional Recurrent Neural Networks for Music Classification." arXiv:1609.04243.  
<https://arxiv.org/abs/1609.04243>

## Source Separation



# Other applications (*discriminative*)



## Music Recommendation

Using deep learning to create personalized music recommendations

Analyze users' preference and listening history.

Well-known *cold start* problem

Aaron van den Oord et al. "Deep Content-based Music Recommendation" (2013)

<https://proceedings.neurips.cc/paper/2013/file/b3ba8f1bee1238a2f37603d90b58898d-Paper.pdf>

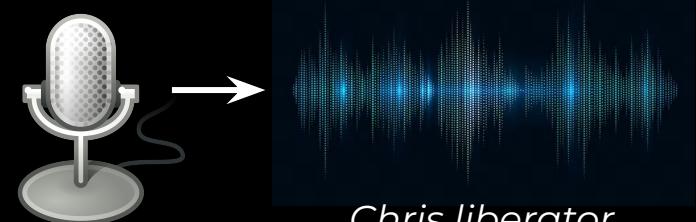
## Query-by-Humming

Fingerprinting for efficient music retrieval in large databases.

Mapping the query and songs to a common reduced space

Provides efficient query-by-humming system

Mostafa, N. et al (2017). A Note Based Query By Humming System Using Convolutional Neural Network. In Interspeech  
[https://www.isca-speech.org/archive\\_v0/Interspeech\\_2017/pdfs/1590.PDF](https://www.isca-speech.org/archive_v0/Interspeech_2017/pdfs/1590.PDF)



Chris liberator



## Audio-to-Score Conversion

Transcribing audio recordings into sheet music using deep models.

Non-uniqueness of the solution and polyphonic problem

Almost solved for piano but still hard for orchestra

Hawthorne et al. (2018). "Onsets and Frames: Dual-Objective Piano Transcription." ISMIR 2018.  
<https://arxiv.org/abs/1710.11153>

# Other applications (generative)



## Symbolic music generation

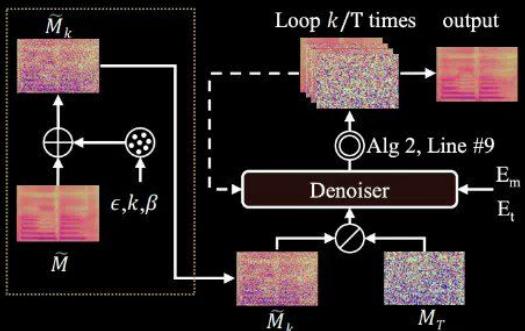
Generating symbolic music with advanced models

Handle long-range dependencies for coherent musical structures.

Huang, C. Z. A., & Du, S. (2020). "Pop Music Transformer: Generating Music with Rhythm and Harmony." *arXiv:2002.00212*.  
<https://arxiv.org/abs/2002.00212>

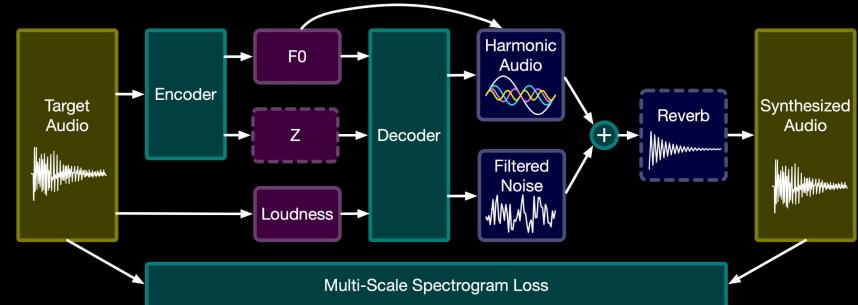
Generating novel synthesized instruments and soundscapes  
Current trends increasingly mixes traditional DSP elements

Engel, J. et al. (2020). "DDSP: Differentiable Digital Signal Processing." *ICLR. 2020*  
<https://openreview.net/forum?id=B1x1ma4tDr>



Analyze structure to create seamless transitions between songs  
Enables automated DJing and playlist creation.

Huang, J. et al. Modeling the compatibility of stem tracks to generate music mashups. In *AAAI 2021*  
<https://ojs.aaai.org/index.php/AAAI/article/view/16092/15899>



## Text-to-singing

Singing voice generation based on input text

Also needs to handle prosody and emotional response

Liu, J. et al. (2022). *Diffsinger: Singing voice synthesis via shallow diffusion mechanism*. In *AAAI 2022*.  
<https://ojs.aaai.org/index.php/AAAI/article/view/21350/21099>

## Algorithmic DJing

